CompSci 471

Artificial Intelligence

Lesson 3: LLM Behavior Investigation Workshop

LLM-as-system component

Mission brief → LLM output

↓

Trust audit

Investigate → Compare → Audit → Brief

input	response	excerpt	judgment	report
Prompt	Output	Evidence	Trust	Brief

Use an LLM as the experimental platform, not as the whole definition of AI.

DEPARTMENT OF COMPUTER AND CYBER SCIENCES • UNITED STATES AIR FORCE ACADEMY

Fall 2026

Workshop frame

Today is a guided investigation, not a lecture.

1. **Explain how LLMs are one type of AI system, with characteristic strengths and weaknesses**
2. **Evaluate how changes in prompts affect LLM outputs**
3. **Identify inconsistencies and failure cases in LLM-based AI systems**
4. **Explain how input design influences output quality across AI systems**

Workshop cycle

Predict → Change → Observe → Compare → Audit → Brief

Operation Lifeline

Investigate different responses to a mission.

Mission
An autonomous aircraft must deliver emergency medical supplies to an isolated field team after a severe storm.

Route North
18 km
Strong headwinds
Near restricted airspace
42% normal-weather energy
16 minutes

Route South
27 km
Lower winds
No known restriction
48% normal-weather energy
23 minutes

Mission conditions
Battery 62%; reserve 15%
Deadline 25 minutes
Communications may be intermittent

Intentional gap
Wind-adjusted energy consumption is not provided.

Authority
A human operator must authorize the final route.

Baseline prompt: “Recommend the best route for the medical-delivery mission. Make one clear recommendation.”

Workshop schedule

Groups investigate for about 20 minutes, then brief findings.

Time	Activity	Purpose
0–5 min	Briefing	Objectives, scenario, rules
5–8 min	Groups + tracks	Assign roles and investigation focus
8–29 min	Investigation	Run trials and collect evidence
29–32 min	Brief prep	Prepare 2–3 minute finding
32–51 min	Presentations	Short group reports
51–53 min	Synthesis	Instructor wrap-up and exit response

3–4 students per group
Assign roles before opening the LLM.

2–3 minute reports
Evidence first; no slides required.

Compare across class
Different prompts or LLMs are encouraged.

Assign roles before opening the LLM

Every group needs an operator, experiment designer, auditor, and briefer.

Operator
Enters prompts. Saves exact outputs. Notes model name and settings.

Experiment designer
Keeps one variable changing at a time. Protects the comparison.

Auditor
Checks assumptions, omissions, contradictions, and unsupported claims.

Recorder / Briefer
Completes the handout and presents the finding.

In a group of three, combine Recorder and Briefer. Rotate roles after the baseline trial if needed.

Experimental rules

Good workshop findings come from controlled comparisons.

1	**Record the system**	LLM/platform, model name if visible, fresh chat or continued chat, and obvious tools/settings.
2	**Change one variable**	Prompt comparison: keep LLM fixed. Model comparison: keep exact prompt fixed.
3	**Preserve evidence**	Save the exact prompt and one short output excerpt that supports your claim.
4	**Avoid sensitive content**	Use only the fictional mission. Do not enter personal, operational, controlled, or assessment information.

A finding must be supported by evidence, not just "the answer seemed better."

Investigation tracks

Each group tests one behavior while the class builds a broader picture.

- **A** Context + constraints
- **B** Missing information
- **C** Ambiguity + wording
- **D** Conflicting facts
- **E** Output structure
- **F** Repeatability / model comparison
- **G** False premise + authority
- **H** Verification + safeguards

If tracks repeat
Duplicate a track with a different approved LLM or a different prompt variant.

If class is large
Use A–F first, then add G–H or split model-comparison groups.

Start with the baseline trial

Everyone begins from the same prompt before changing anything.

Baseline prompt
"Recommend the best route for the medical-delivery mission. Make one clear recommendation."

1. **Predict what may go wrong**
2. **Run the exact prompt**
3. **Record one strength and one weakness**
4. **Identify one thing to verify**

5 min

Do not judge the LLM by polish. Judge by evidence.

Group investigation time

Run controlled trials and preserve evidence.

21 min **Your group should complete at least three trials:**

1. Baseline
Use the exact shared prompt. Record the output.

2. Controlled change
Change one prompt feature or one model.

3. Follow-up trial
Test whether the observed behavior repeats or improves.

Trial	LLM/model	Exact change	Prediction	Key output evidence
Baseline		None		
Trial 2				
Trial 3				

A useful conclusion names the variable changed and cites a short output excerpt.

Trust audit

Classify claims before deciding whether the output should be used.

Classification	Meaning
Supported	Directly grounded in the mission brief
Reasonable inference	Logically supported, but not stated
Unsupported	Added without evidence
Contradicted	Conflicts with provided information
Unverifiable	Cannot be checked with available information

Hidden assumption
claim depends on something not stated

Missed constraint
ignores battery, reserve, deadline, or authority

Unjustified confidence
sounds certain despite missing information

Requested clarification
recognizes data needed before decision

Trust options: usable as written | human review | more information needed | should not be used

Prepare your findings brief

You have 3 minutes to convert evidence into a 2–3 minute report.

1. Investigation
What did you test?

2. Comparison
What changed and what stayed constant?

3. Evidence
What output excerpt supports your finding?

4. Safeguard
What should a user or AI system do next?

3 min

No slide deck required. Present from your final handout page or one screenshot plus notes.

Presentation round

Each group gives a concise evidence-based finding.

Time	Part	What to say
20 sec	Investigation	“We tested whether…”
60 sec	Comparison	Baseline versus changed prompt or model
45 sec	Evidence-based finding	One observed behavior + short output evidence
30 sec	Implication / safeguard	What a user or larger AI system should do

4–6 groups
2.5–3 minutes each

7–8 groups
2 minutes each

Audience task
Record findings from two other groups

Stop at time. The goal is class comparison, not long reports.

Listen for patterns

The class is comparing behaviors across prompts, groups, and LLMs.

Group	Variable tested	Important observed behavior	General lesson

Prompt-specific?
Did wording or output format seem to drive the difference?

Model-specific?
Did a difference appear tied to a particular LLM or repeated run?

General principle?
Did the pattern seem likely to apply across AI systems?

After two presentations, start looking for patterns—not just isolated examples.

Synthesis: what did we learn?

LLM outputs depend on inputs, model behavior, and external verification.

1. **LLMs can summarize, structure, reason, and expose questions quickly**
2. **They may miss constraints, assume missing facts, or sound overconfident**
3. **Prompt design changes what the model notices, assumes, and communicates**
4. **External data, deterministic checks, monitoring, and human review remain essential**

Better prompting can improve communication, but it cannot create missing facts or guarantee truth.

Exit response

Write individually before leaving.

1. What prompt or model change produced the most important output change?
2. What inconsistency, failure, or limitation did your group observe?
3. What should a responsible user or larger AI system verify before acting?
4. What did this workshop reveal about LLMs as AI-system components?

Next class: Problem Formulation — turning real-world tasks into searchable models.